Machine Learning
Recommendation

Jeff Abrahamson

août 2019

Recommendation

Definition

Given data about a user, his environment, and some items of interest (*training data*), determine items to recommend.

Definition

> Given data about a user, his environment, and some items of interest (*training data*), determine items to recommend.

We don't have to find the max k.

It's enough to find k within some max n.

Examples

- Amazon
- Google News (or Le Monde)
- Facebook
- Medical testing
- App Store / Google Play
- Youtube
- Advertising
- Netflix, last.fm, Spotify, Pandora, . . .
- Browser (URL recommendations)
- Search

Client Value Proposition

- Find opportunities
- Reduce choice
- Explore options
- Discover long tails
- Recreation

Provider Value Proposition

- Offer a unique or additional service (beyond competitors)
- Customer trust and loyalty
- Increase sales, CTR, conversions
- Better understand customers

Recommendation

Content-based filtering *(filtrage basée sur le contenu)*	More things similar to what I like
Collaborative filtering *(filtrage collaboratif)*	More of what other people who like what I like like
Knowledge-based filtering *(filtrage basée sur connaissance)*	More of what I need.

Content-based filtering

More things similar to what I like
Plus de ce qui ressemble à ce que j'aime

Advantages

- No need for community
- Possible to compare items

Disadvantages

- Understand content
- Cold start problem
- Serendipity

Collaborative filtering

More of what other people who like what I like like
Plus de ce que d'autres qui aiment ce que j'aime aiment

Advantages

- need to understand content
- Serendipity
- Learn market

Disadvantages

- User feedback
- Cold start problem (users)
- Cold start problem (items)

Knowledge-based filtering

More of what I need
Plus de ce qu'il faut

Advantages

- Deterministic
- Certainty
- Cold start problem
- Market knowledge

Disadvantages

- Studies to bootstrap
- Static model, doesn't learn from trends

Utility Matrix

- Users (utilisateurs)
- Items (objets)

Utility Matrix

- Users (utilisateurs)
- Items (objets)

The goal is to fill in the blanks.

	I_1	I_2	I_3	I_4	I_5
U_1	1				
U_2			1	1	1
U_3		1		1	1

Example: books sales at Amazon.

But thousands or millions of columns and rows.

Utility Matrix

- Users (utilisateurs)
- Items (objets)

The goal is to fill in the blanks.

	I_1	I_2	I_3	I_4	I_5
U_1	3				
U_2			5	1	4
U_3		2		5	1

Example: film advice at Netflix.

But thousands or millions of columns and rows.

Utility Matrix

How do we make the matrix?

- Ask users
- Observe users

That's usually expensive. . .

Item Profiles

Examples:

- Films $\Rightarrow$?
- Books $\Rightarrow$?
- News $\Rightarrow$?
- Images $\Rightarrow$?

Item Profiles

Examples:

- Films $\Rightarrow$?
- Books $\Rightarrow$?
- News $\Rightarrow$?
- Images $\Rightarrow$?

Films :

Content: acters, directors, year (decade, etc.), length

Collaborative: seen, opinion (1–5), when seen relative to release

Item Profiles

Examples:

- Films   ⇒ ?
- Books   ⇒ ?
- News   ⇒ ?
- Images   ⇒ ?

Books:

Content : authers, genre, year (decade, etc.), number of pages, content (very difficult)

Collaborative: read, opinion (1–5), how read

Item Profiles

Examples:

- Films $\Rightarrow$?
- Books $\Rightarrow$?
- News $\Rightarrow$?
- Images $\Rightarrow$?

News:

Content : source, section, TF-IDF word vectors

Collaborative:

Item Profiles

Examples:

- Films $\Rightarrow$?
- Books $\Rightarrow$?
- News $\Rightarrow$?
- Images $\Rightarrow$?

Images :

Content:

Collaborative:

Item Profiles

Examples:

- Films $\Rightarrow$?
- Books $\Rightarrow$?
- News $\Rightarrow$?
- Images $\Rightarrow$?

Also: user profile, user behavior

Mathematics

Vectors

Similarity

Similarity : Jaccard Index

or: *Indice de Jaccard*, *Jaccard similarity coefficient*

Similarity:

$$J(A, B) = \frac{|A \cap B|}{|A \cup B|}$$

Similarity : Jaccard Index

Similarity:

$$J(A, B) = \frac{|A \cap B|}{|A \cup B|}$$

Distance:

$$J_\delta(A, B) = 1 - J(A, B)$$

cosine similarity

or: *mesure cosinus*, *similarité cosinus*

Similarity:

$$\cos \theta = \frac{A \cdot B}{\| A \| \| B \|}$$

cosine similarity

Similarity:

$$S_C(A, B) = \frac{A \cdot B}{\| A \| \| B \|}$$

cosine similarity

Similarity:

$$S_C(A, B) = \frac{A \cdot B}{\| A \| \| B \|}$$

Distance:

$$D_C(A, B) = 1 - S_C(A, B)$$

cosine similarity

Similarity:

$$S_C(A, B) = \frac{A \cdot B}{\| A \| \| B \|}$$

Distance:

$$D_C(A, B) = 1 - S_C(A, B)$$

We only consider non-empty components in the vector.

Text: an application of cosine similarity

Texts: TF-IDF

- Vectors of word frequencies
- Frequency $\not\Rightarrow$ significance

Texts: TF-IDF

- Vectors of word frequencies
- Frequency $\not\Rightarrow$ significance
- Term Frequency - Inverse Document Frequency

Texts: TF-IDF

$$TF_{ij} = \frac{f_{ij}}{\max_k f_{kj}} \qquad IDF_I = \log_2\left(\frac{N}{n_i}\right)$$

$$TF\text{-}IDF_{ij} = TF_{ij} \cdot IDF_i$$

with :

$f_{ij} =$ frequency of word i in document j

$N =$ number of documents

$n_i =$ number of documents in which we find word i

Texts: TF-IDF

$$TF_{ij} = \frac{f_{ij}}{\max_k f_{kj}} \qquad IDF_I = \log_2 \left(\frac{N}{n_i} \right)$$

$$TF\text{-}IDF_{ij} = TF_{ij} \cdot IDF_i$$

with :

f_{ij} = frequency of word i in document j

N = number of documents

n_i = number of documents in which we find word i

IDF is a mesure of how much information a word carries

TF-IDF tells us which words best characterise a document

Texts: TF-IDF

$$TF_{ij} = \frac{f_{ij}}{\max_k f_{kj}} \qquad IDF_l = \log_2\left(\frac{N}{n_i}\right)$$

$$TF\text{-}IDF_{ij} = TF_{ij} \cdot IDF_i$$

with :

f_{ij} = frequency of word i in document j

N = number of documents

n_i = number of documents in which we find word i

IDF is a mesure of how much information a word carries

TF-IDF tells us which words best characterise a document

Variation: boolean, log, stop word filtering

Content-Based Filtering

	I_1	I_2	I_3	I_4	I_5	I_6	I_7
U_1	3				2	4	
U_2			5	1	4		3
U_3		2		5	1	5	

More things similar to what I like
Plus de ce qui ressemble à ce que j'aime

	I_1	I_2	I_3	I_4	I_5	I_6	I_7
U_1	3				2	4	
U_2			5	1	4		3
U_3		2		5	1	5	

More things similar to what I like
Plus de ce qui ressemble à ce que j'aime

Then, we can cluster (*regroupement*, *partitionnement de données*), etc.

Content-Based Filtering

Based on item profiles

- More stable (in principle)
- $O(n^2)$ (but often less, items often aren't categorised together)
- Can reduce to threshold
- Can pre-calculate, queries become faster

Collaborative Filtering

	I_1	I_2	I_3	I_4	I_5
U_1	3		4	2	
U_2			5	1	4
U_3		2		5	1

More of what other people who like what I like like
Plus de ce que d'autres qui aiment ce que j'aime aiment

Collaborative Filtering

	I_1	I_2	I_3	I_4	I_5
U_1	3				
U_2			5	1	4
U_3		2		5	1

User profile

Collaborative Filtering

	I_1	I_2	I_3	I_4	I_5
U_1	3		4	2	
U_2			5	1	4
U_3		2		5	1

Item profile

Utility Matrix Symmetry

- Propose items based on users
- Propose users based on items

Utility Matrix Symmetry

- Propose items based on users
- Propose users based on items

But remember: 2 items being similar $\not\Rightarrow$ 2 users similar.

Utility Matrix Symmetry

- Propose items based on users
- Propose users based on items

But remember: 2 items being similar $\not\Rightarrow$ 2 users similar.

Thought experiment: consider comparing people vs comparing objects.

Utility Matrix Symmetry

- Propose items based on users
- Propose users based on items

To estimate $m_{u,i}$,

- Find k users like U_u
- Find k items like I_i

Utility Matrix : Estimate $m_{u,i}$ (naive averaging)

	I_1	I_2	I_3	I_4	I_5
U_1	3		4	2	□
U_2			5	1	4
U_3		2		2	3

- Find k users like U_u, take $\frac{1}{k} \sum_{j=1}^{k} m_{u_j,i}$
- Find k items like I_i, take $\frac{1}{k} \sum_{j=1}^{k} m_{u,i_j}$

Utility Matrix : Estimate $m_{u,i}$ (naive averaging)

	I_1	I_2	I_3	I_4	I_5
U_1	3		4	2	$\square$
U_2			5	1	4
U_3		2		2	3

- Find k users like U_u, take $\frac{1}{k} \sum_{j=1}^{k} m_{u_j,i}$
- Find k items like I_i, take $\frac{1}{k} \sum_{j=1}^{k} m_{u,i_j}$

We have to compute the entire line (or the part which is likely to be important)

Utility Matrix : Estimate $m_{u,i}$ (naive averaging)

	I_1	I_2	I_3	I_4	I_5
U_1	3		4	2	□
U_2			5	1	4
U_3		2		2	3

- Find k users like U_u, take $\frac{1}{k} \sum_{j=1}^{k} m_{u_j,i}$
- Find k items like I_i, take $\frac{1}{k} \sum_{j=1}^{k} m_{u,i_j}$

Once we've computed U_u, the other k users lets us take a shortcut.

Utility Matrix : Estimate $m_{u,i}$ (naive averaging)

	I_1	I_2	I_3	I_4	I_5
U_1	3		4	2	☐
U_2			5	1	4
U_3		2		2	3

- Find k users like U_u, take $\frac{1}{k} \sum_{j=1}^{k} m_{u_j,i}$
- Find k items like I_i, take $\frac{1}{k} \sum_{j=1}^{k} m_{u,i_j}$

For I_i, we have to compute most of the I_j before we can fill in a single line. But item-item filters are often more reliable.

Utility Matrix : Estimate $m_{u,i}$ (naive averaging)

	I_1	I_2	I_3	I_4	I_5
U_1	3		4	2	$\square$
U_2			5	1	4
U_3		2		2	3

- Find k users like U_u, take $\frac{1}{k} \sum_{j=1}^{k} m_{u_j,i}$
- Find k items like I_i, take $\frac{1}{k} \sum_{j=1}^{k} m_{u,i_j}$

In any case, we can mostly precompute in advance.

Utility Matrix

The matrix is sparse.

$\implies$ clustering $\implies$ reduced matrix

Utility Matrix

The matrix is sparse.

$\implies$ clustering $\implies$ reduced matrix

Estimate on the reduced matrix, then take items and users as representative for the cluster.

Amazon : Item-to-Item Collaborative Filtering

Observations :

<div style="color: blue; text-align: center;">Clustering is expensive, reduces quality</div>

Amazon : Item-to-Item Collaborative Filtering

Observations :

Dimension reduction reduces quality

Amazon : Item-to-Item Collaborative Filtering

Observations :

Users interact with very few items

Amazon : Item-to-Item Collaborative Filtering

Observations :

Rapid response desirable

Amazon : Item-to-Item Collaborative Filtering

Scales independent of the number of users or of items

- Online
- Offline

G. Linden, B. Smith, J. York, *Amazon.com Recommendations: Item-to-Item Collaborative Filtering*, Internet Computing (7, 1), 22 Jan 2003.

Amazon : Item-to-Item Collaborative Filtering

Offline (Precomputation)

for *each item I_1 to sell* **do**
    for *each user C who has purchased I_1* **do**
        for *each item I_2 bought by C* **do**
            (I_1, I_2)++
        end
    end
    for *each item I_2* **do**
        $S_{I_1, I_2} \leftarrow S(I_1, I_2)$
    end
end

Slope One

Linear regression on user opinions (ratings)

Daniel Lemire and Anna Maclachlan, *Slope One Predictors for Online Rating-Based Collaborative Filtering*,Proceedings of SIAM Data Mining (SDM) 2005.

Slope One : algorithm

Offline :

for *chaque I_i, I_j* **do**

$\quad \mathcal{U} \leftarrow \{$users who have expressed an opinion on $I_i, I_j\}$

$\quad \text{dev}_{i,j} \leftarrow \dfrac{1}{\parallel \mathcal{U} \parallel} \sum\limits_{u \in \mathcal{U}} (r_u(i) - r_u(j))$

end

Online (for u) :

$\mathcal{V} \leftarrow \{j \mid u \text{ has expressed an opinion on } I_j\}$

$r_u(i) \leftarrow \dfrac{1}{\parallel \mathcal{V} \parallel} \sum\limits_{u \in \mathcal{V}} (\text{dev}_{i,j} - r_u(j))$

Dimensionality Reduction

Dimensionality reduction

SVD, typically $k = 20 \ldots 100$

$$M = U\Sigma V^*$$

Dimensionality reduction

SVD, typically $k = 20 \ldots 100$

$$\begin{pmatrix} a_1 & \cdots & a_m \end{pmatrix} \begin{pmatrix} b_1 \\ \vdots \\ b_n \end{pmatrix} = \text{scalar}$$

Dimensionality reduction

SVD, typically $k = 20 \ldots 100$

$$\begin{pmatrix} a_1 \\ \vdots \\ a_m \end{pmatrix} \begin{pmatrix} b_1 & \cdots & b_n \end{pmatrix} = \begin{pmatrix} c_{1,1} & \cdots & c_{1,n} \\ \vdots & & \vdots \\ c_{m,1} & \cdots & c_{m,n} \end{pmatrix}$$

Dimensionality reduction

SVD, typically $k = 20 \ldots 100$

$$\begin{pmatrix} a_{1,1} & a_{1,2} & a_{1,3} \\ \vdots & \vdots & \vdots \\ a_{m,1} & a_{m,2} & a_{m,3} \end{pmatrix} \begin{pmatrix} b_{1,1} & \cdots & b_{1,n} \\ b_{2,1} & \cdots & b_{2,n} \\ b_{3,1} & \cdots & b_{3,n} \end{pmatrix} = \begin{pmatrix} c_{1,1} & \cdots & c_{1,n} \\ \vdots & & \vdots \\ c_{m,1} & \cdots & c_{m,n} \end{pmatrix}$$

Dimensionality reduction

SVD, typically $k = 20 \ldots 100$

$$\begin{pmatrix} a_{1,1} & \cdots & a_{1,k} \\ \vdots & & \vdots \\ a_{m,1} & \cdots & a_{m,k} \end{pmatrix} \begin{pmatrix} c_{1,1} & \cdots & c_{1,n} \\ \vdots & & \vdots \\ c_{k,1} & \cdots & c_{k,n} \end{pmatrix} = \begin{pmatrix} c_{1,1} & \cdots & c_{1,n} \\ \vdots & & \vdots \\ c_{m,1} & \cdots & c_{m,n} \end{pmatrix}$$

Challenges

- How do we measure success?
- What are our features?

Campus Académie Ouest, août 2019

Clustering

- kNN
- Curse of Dimensionality
- Scalability

Clustering

- kNN *k*-Nearest Neighbor
- Curse of Dimensionality
- Scalability 10^7 clients, 10^6 objets